

Name: Syed Rafay Hassan

Roll No: 22i-1955

## 1. Network Details

We tested two deep learning models for emotion recognition: **ResNet50** and **EfficientNetB0**.

- Both models were first trained on ImageNet and then used as the base.
- On top of them, we added two heads:
  - **Classification head:** Dropout → Dense layer (256 units, ReLU) → Dense (8 classes, Softmax).
  - **Regression head:** Dropout → Dense layer (128 units, ReLU) → Dense (2 values, Tanh).
- **Optimizer:** Adam (learning rate = 0.0001).
- **Losses:** Categorical cross-entropy (for classification) and Mean Squared Error (for valence/arousal regression). Both were given equal weight.
- **Training:** 12 epochs, batch size = 32, input size = 224×224.

## 2. Dataset Splits

- Dataset: **DL\_Assignment1\_Dataset**
- **Task 1 (Classification):** Predict 8 emotion categories.
- **Task 2 (Regression):** Predict valence and arousal values.
- **Splits:** 85% training and 15% validation. The split was stratified to keep class balance.

## 3. Training Curves

- **ResNet50:** Accuracy increased slowly and stabilized around **40% validation accuracy**.
- **EfficientNetB0:** Learned faster but plateaued at a lower **28% validation accuracy**.
- Both models showed steady loss reduction without major overfitting.

## 4. Results on Validation Set

### Classification Results

## Model Accuracy F1-Macro Mean AP

**ResNet50** **0.408** **0.405** 0.396

EfficientNetB0 0.278 0.259 0.335

- ResNet50 performed much better in classification, especially in terms of F1-score.
- Confusion matrices showed that EfficientNetB0 often confused high-frequency classes.

## Regression Results

Model	RMSE (Valence)	RMSE (Arousal)	Corr (Valence)	Corr (Arousal)	CCC (Valence)	CCC (Arousal)	SAGR (Valence)	SAGR (Arousal)
<b>ResNet50</b>	<b>0.425</b>	<b>0.361</b>	<b>0.484</b>	<b>0.388</b>	<b>0.434</b>	<b>0.320</b>	<b>0.743</b>	<b>0.778</b>
EfficientNetB0	0.429	0.384	0.441	0.244	0.323	0.187	0.712	0.765

- ResNet50 gave better results for both valence and arousal across all metrics.
- EfficientNetB0 performed weaker, especially in correlation and CCC scores.

## 5. Comparison and Observations

- **ResNet50 was the best overall model.** It achieved higher accuracy, F1-score, and better regression results compared to EfficientNetB0.
- **EfficientNetB0 underperformed** and could not match ResNet50's balance between classification and regression.
- The results suggest that **ResNet50 is more reliable for emotion recognition tasks** on this dataset.